

Personalized and Private Peer-to-Peer Machine Learning

(novembre 2018)


Pierre-Louis Biojout, Leonid Kniazev

Learning from connected devices data

- Connected devices are widespread and collect increasingly large and sensitive user data
- Ex: browsing logs, health data, accelerometer, geolocation...
- Great opportunity for providing personalized services but raises
- serious privacy concerns
- Centralize data from all devices: best for utility, bad for privacy
- Learn on each device separately: best for privacy, bad for utility

Proposed solution: fully decentralized network

- Personal data stays on user's device
- Peer-to-peer and asynchronous communication
- No single point of failure/entry as in server-client architecture
- Scalability-by-design to many devices through local updates (see e.g. NIPS 2017 paper [Lian et al., 2017])
- Learn a personalized model for each user (multi-task learning)
- General idea: trade-off between model accuracy on local data and smoothness with respect to similar users



Problem setting

Problem setting


- A set $V = \llbracket n \rrbracket = \{1, \dots, n\}$ of n learning agents
- A **convex** loss function $\ell : \mathbb{R}^p \times \mathcal{X} \times \mathcal{Y}$
- **Personalized and imbalanced data**: agent i has dataset $\mathcal{S}_i = \{(x_i^j, y_i^j)\}_{j=1}^{m_i}$ of size $m_i \geq 0$ drawn from μ_i
- **Purely local model**: agent i can learn a model θ_i on its own by minimizing the loss on its local data

$$\mathcal{L}_i(\theta) = \frac{1}{m_i} \sum_{j=1}^{m_i} \ell(\theta; x_i^j, y_i^j) + \lambda_i \|\theta\|^2, \text{ with } \lambda_i \geq 0$$

Problem setting

- Network: weighted connected graph $G = (V, E)$
- $E \subseteq V \times V$ set of undirected edges
- Weight matrix $W \in \mathbb{R}^{n \times n}$: symmetric, nonnegative, with $W_{ij} = 0$ if $(i, j) \notin E$ or $i = j$
- **Assumption:** network weights are given and represent the underlying similarity between agents

Problem setting



- Agents have only a **local view** of the network: they only know their neighborhood $\mathcal{N}_i = \{j \neq i : W_{ij} > 0\}$ and associated weights

Problem formulation

- Denoting $\Theta = [\Theta_1; \dots; \Theta_n] \in \mathbb{R}^{np}$, we use a graph regularization formulation [Evgeniou and Pontil, 2004, Vanhaesebrouck et al., 2017]:

$$\min_{\Theta \in \mathbb{R}^{np}} \mathcal{Q}_{\mathcal{L}}(\Theta) = \frac{1}{2} \sum_{i < j}^n w_{ij} \|\Theta_i - \Theta_j\|^2 + \mu \sum_{i=1}^n D_{ii} c_i \mathcal{L}_i(\Theta_i; \mathcal{S}_i)$$

- $\mu > 0$ trade-off parameter, $D_{ii} = \sum_j w_{ij}$ normalization factor
- $c_i \in (0, 1] \propto m_i$ is the “confidence” of agent i
- Implements a trade-off between having similar models for strongly connected agents and models that are accurate on their respective local datasets

Non-private decentralized algorithm

Non-private decentralized algorithm

- Time and communication models:
 - **Asynchronous time**: each agent has a random local clock and wakes up when it ticks
 - **Broadcast communication**: agents send messages to all their neighbors at once (without expecting a reply)
- **Algorithm**: assume agent i wakes up at step t
 1. Agent i updates its model based on information from neighbors:

$$\Theta_i(t+1) = (1 - \alpha)\Theta_i(t) + \alpha \left(\sum_{j \in \mathcal{N}_i} \frac{W_{ij}}{D_{ii}} \Theta_j(t) - \mu c_i \nabla \mathcal{L}_i(\Theta_i(t); \mathcal{S}_i) \right)$$

2. Agent i sends its updated model $\Theta_i(t+1)$ to its neighborhood \mathcal{N}_i

Private algorithm

Private algorithm

- In some applications, **data may be sensitive** and agents may not want to reveal it to anyone else
- In our algorithms, the agents never communicate their local data but **exchange sequences of models computed from data**
- Consider an adversary observing **all the information sent over the network** (but not the internal memory of agents)
- **Goal:** how can we guarantee that no/little information about the local dataset is leaked by the algorithm?

Differential privacy

(ϵ, δ) -Differential Privacy

Let \mathcal{M} be a randomized mechanism taking a dataset as input, and let $\epsilon > 0, \delta \geq 0$. We say that \mathcal{M} is (ϵ, δ) -differentially private if for all datasets $\mathcal{S}, \mathcal{S}'$ differing in a single data point and for all sets of possible outputs $\mathcal{O} \subseteq \text{range}(\mathcal{M})$, we have:

$$\Pr(\mathcal{M}(\mathcal{S}) \in \mathcal{O}) \leq e^\epsilon \Pr(\mathcal{M}(\mathcal{S}') \in \mathcal{O}) + \delta.$$

- Guarantees that \mathcal{M} does not leak much information about any individual data point
- Information-theoretic (no computational assumptions)

Differentially private algorithm

- Differentially-private algorithm:

1. Replace the update of the algorithm by

$$\tilde{\Theta}_i(t+1) = (1-\alpha)\tilde{\Theta}_i(t) + \alpha \left(\sum_{j \in \mathcal{N}_i} \frac{W_{ij}}{D_{ii}} \tilde{\Theta}_j(t) - \mu c_i (\nabla \mathcal{L}_i(\tilde{\Theta}_i(t); \mathcal{S}_i) + \eta_i(t)) \right),$$

where $\eta_i(t) \sim \text{Laplace}(0, s_i(t))^p \in \mathbb{R}^p$

2. Agent i then broadcasts noisy iterate $\tilde{\Theta}_i(t+1)$ to its neighbors

Privacy guarantee

Theorem ([Bellet et al., 2017])

Let $i \in \llbracket n \rrbracket$ and assume

- $\ell(\cdot; x, y)$ L_0 -Lipschitz w.r.t. the L_1 -norm for all (x, y)
- Agent i wakes up on iterations $t_i^1, \dots, t_i^{T_i}$
- For some $\epsilon_i(t_i^k) > 0$, the noise scale is $s_i(t_i^k) = \frac{2L_0}{\epsilon_i(t_i^k)m_i}$

Then for any initial point $\tilde{\Theta}(0)$ independent of \mathcal{S}_i , the mechanism $\mathcal{M}_i(\mathcal{S}_i)$ is $(\bar{\epsilon}_i, 0)$ -DP with $\bar{\epsilon}_i = \sum_{k=1}^{T_i} \epsilon_i(t_i^k)$.

- **Sweet spot**: the less data, the more noise added by the agent, but the least influence in the network

Privacy/utility trade-off

Theorem ([Bellet et al., 2017])

For any $T > 0$, let $(\tilde{\Theta}(t))_{t=1}^T$ be the sequence of iterates generated by T iterations. Under appropriate assumptions, we have:

$$\mathbb{E} [\mathcal{Q}_{CL}(\Theta(T)) - \mathcal{Q}_{CL}^*] \leq (1 - \rho)^T (\mathcal{Q}_{CL}(\Theta(0)) - \mathcal{Q}_{CL}^*) + \text{additive error}.$$

- **Second term** gives additive error due to noise
- More results in the paper

Experiments

Experiments

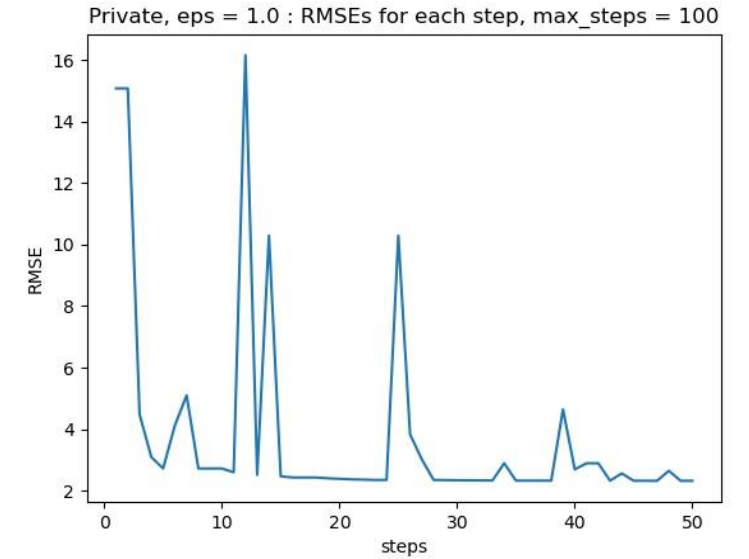
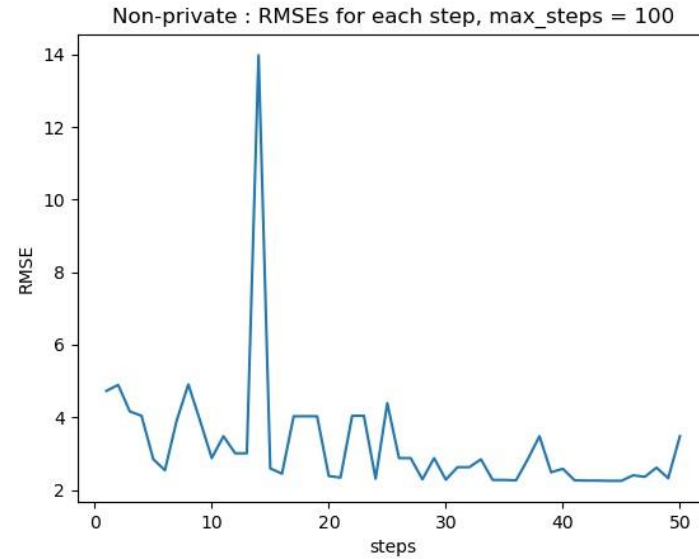
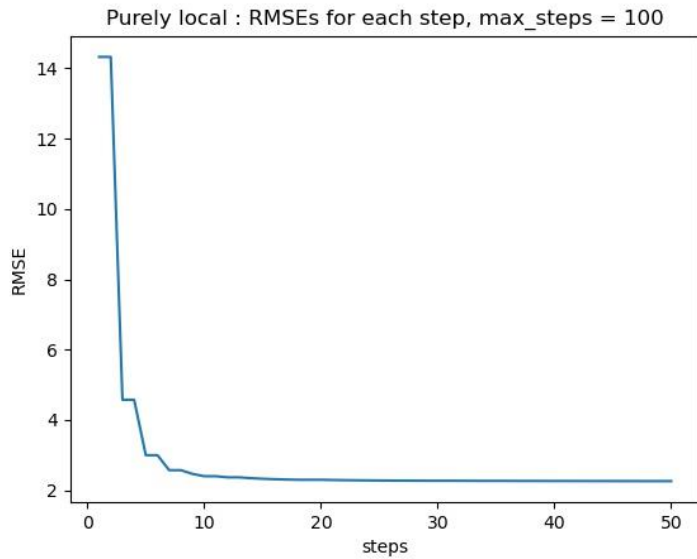
- Implementation in Python
- Simulation of the network on a single computer
- Use of ml-100k dataset
- The network is obtained by setting $W_{ij} = 1$ if agent i is within the 10-nearest neighbors of agent j

Results

	Purely local models	Non-priv. CD	Priv. $\bar{\epsilon} = 1$	Priv. $\bar{\epsilon} = 0.5$	Priv. $\bar{\epsilon} = 0.1$
Published results	1.2834	0.9502	0.9527	0.9545	0.9855
Our results	2.5821	2.5676	2.6524	2.4766	2.9688

Per-user test RMSE (averaged over users and 5 random runs) on MovieLens-100K.

Graphs



Conclusion

Acknowledgement

- Based on the paper *Personalized and Private Peer-to-Peer Machine Learning* (Aurélien Bellet and al.)
- Part of the illustrations and formulas in our slides are from the slides *personalized and private peer-to-peer machine learning* (NIPS 2017 workshop on “Machine Learning on the Phone and other Consumer Devices” Long Beach, December 9, 2017)